

Active Burn Forecast

- User_ID: (used only for tracking individuals, non-predictive)
- Gender: (used to capture differences in calorie burn between males and females)
- Age: (age influences metabolism and calorie burn rates)
- Height: (taller individuals may burn more calories)
- Weight: (heavier individuals generally burn more calories during activity)
- Duration: (longer exercise duration typically results in more calories burned)
- Heart_Rate: (higher heart rates indicate more intense activity, leading to higher calorie burn)
- Body_Temp: (body temperature can correlate with the intensity of exercise and calorie burn)
- Calories: (the target variable representing the number of calories burned during the activity, which the model aims to predict)

Step 1 : import libraries & create dataframe

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn.linear_model import LinearRegression
#from sklearn.linear_model import Ridge,Lasso
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
from statsmodels.stats.outliers_influence import
variance_inflation_factor
import pickle

import warnings
from warnings import filterwarnings
filterwarnings("ignore")

sns.set()

#Load the Calories dataset
df1 = pd.read_csv("calories.csv")
df1.head()
```

| | User_ID | Calories |
|---|----------|----------|
| 0 | 14733363 | 231.0 |
| 1 | 14861698 | 66.0 |
| 2 | 11179863 | 26.0 |
| 3 | 16180408 | 71.0 |
| 4 | 17771927 | 35.0 |

```
df1.shape
```

```
(15000, 2)
```

```
#Load the Exercise Dataset
```

```
df2 = pd.read_csv("exercise.csv")
```

```
df2.head()
```

| | User_ID | Gender | Age | Height | Weight | Duration | Heart_Rate |
|---|----------|--------|-----|--------|--------|----------|------------|
| 0 | 14733363 | male | 68 | 190.0 | 94.0 | 29.0 | 105.0 |
| 1 | 14861698 | female | 20 | 166.0 | 60.0 | 14.0 | 94.0 |
| 2 | 11179863 | male | 69 | 179.0 | 79.0 | 5.0 | 88.0 |
| 3 | 16180408 | female | 34 | 179.0 | 71.0 | 13.0 | 100.0 |
| 4 | 17771927 | female | 27 | 154.0 | 58.0 | 10.0 | 81.0 |

```
df2.shape
```

```
(15000, 8)
```

Now Concatenate both the Dataframe i.e df1 and df2

```
df = pd.concat([df2,df1["Calories"]],axis=1)
```

```
df.head()
```

| | User_ID | Gender | Age | Height | Weight | Duration | Heart_Rate |
|---|----------|--------|-----|--------|--------|----------|------------|
| 0 | 14733363 | male | 68 | 190.0 | 94.0 | 29.0 | 105.0 |
| 1 | 14861698 | female | 20 | 166.0 | 60.0 | 14.0 | 94.0 |
| 2 | 11179863 | male | 69 | 179.0 | 79.0 | 5.0 | 88.0 |
| 3 | 16180408 | female | 34 | 179.0 | 71.0 | 13.0 | 100.0 |
| 4 | 17771927 | female | 27 | 154.0 | 58.0 | 10.0 | 81.0 |

| | Calories |
|---|----------|
| 0 | 231.0 |
| 1 | 66.0 |
| 2 | 26.0 |
| 3 | 71.0 |
| 4 | 35.0 |

Step 2 : Data cleaning (handling the null values and dropping unwanted columns)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User_ID         15000 non-null  int64
1   Gender          15000 non-null  object
2   Age             15000 non-null  int64
3   Height          15000 non-null  float64
4   Weight          15000 non-null  float64
5   Duration        15000 non-null  float64
6   Heart_Rate      15000 non-null  float64
7   Body_Temp       15000 non-null  float64
8   Calories        15000 non-null  float64
dtypes: float64(6), int64(2), object(1)
memory usage: 1.0+ MB
```

```
df.describe()
```

| | User_ID | Age | Height | Weight |
|------------|--------------|--------------|--------------|--------------|
| Duration \ | | | | |
| count | 1.500000e+04 | 15000.000000 | 15000.000000 | 15000.000000 |
| mean | 1.497736e+07 | 42.789800 | 174.465133 | 74.966867 |
| std | 2.872851e+06 | 16.980264 | 14.258114 | 15.035657 |
| min | 1.000116e+07 | 20.000000 | 123.000000 | 36.000000 |
| 25% | 1.247419e+07 | 28.000000 | 164.000000 | 63.000000 |
| 50% | 1.499728e+07 | 39.000000 | 175.000000 | 74.000000 |
| 75% | 1.744928e+07 | 56.000000 | 185.000000 | 87.000000 |
| max | 1.744928e+07 | 56.000000 | 185.000000 | 87.000000 |

```
max    1.999965e+07    79.000000    222.000000    132.000000
30.000000
```

```
      Heart_Rate    Body_Temp    Calories
count  15000.000000  15000.000000  15000.000000
mean    95.518533    40.025453    89.539533
std     9.583328     0.779230    62.456978
min     67.000000    37.100000     1.000000
25%     88.000000    39.600000    35.000000
50%     96.000000    40.200000    79.000000
75%    103.000000    40.600000   138.000000
max    128.000000    41.500000   314.000000
```

```
df.isnull().sum()
```

```
User_ID      0
Gender       0
Age          0
Height       0
Weight       0
Duration     0
Heart_Rate   0
Body_Temp    0
Calories     0
dtype: int64
```

```
# drop User_ID column because this is not required from Main Dataframe itself
```

```
df.drop(columns = ["User_ID"],axis=1,inplace =True)
```

```
df.head()
```

```
      Gender  Age  Height  Weight  Duration  Heart_Rate  Body_Temp
Calories
0    male    68   190.0   94.0     29.0     105.0     40.8
231.0
1  female    20   166.0   60.0     14.0     94.0     40.3
66.0
2    male    69   179.0   79.0      5.0     88.0     38.7
26.0
3  female    34   179.0   71.0     13.0    100.0     40.5
71.0
4  female    27   154.0   58.0     10.0     81.0     39.8
35.0
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15000 entries, 0 to 14999
Data columns (total 8 columns):
```

| # | Column | Non-Null Count | Dtype |
|---|------------|----------------|---------|
| 0 | Gender | 15000 non-null | object |
| 1 | Age | 15000 non-null | int64 |
| 2 | Height | 15000 non-null | float64 |
| 3 | Weight | 15000 non-null | float64 |
| 4 | Duration | 15000 non-null | float64 |
| 5 | Heart_Rate | 15000 non-null | float64 |
| 6 | Body_Temp | 15000 non-null | float64 |
| 7 | Calories | 15000 non-null | float64 |

dtypes: float64(6), int64(1), object(1)
memory usage: 937.6+ KB

step 3 : Encoding

Separate Categorical and Numerical Features

1. Categorical Feature


```
#Fatching Categorical Data
cat_col=[col for col in df.columns if df[col].dtype=='O'] #-
>Object-"o"
cat_col

['Gender']

df["Gender"].value_counts()

Gender
female    7553
male      7447
Name: count, dtype: int64

# plotting the gender column in count plot
sns.countplot(df['Gender'])
plt.show()
```



```
categorical = df[cat_col]
categorical.head()
```

```
  Gender
0   male
1  female
2   male
3  female
4  female
```

```
categorical = pd.get_dummies(categorical["Gender"],drop_first=True)
```

```
categorical
```

```
   male
0    True
1   False
2    True
3   False
4   False
...    ...
14995 False
14996 False
14997 False
14998  True
```

```
14999    True
```

```
[15000 rows x 1 columns]
```

2.Numerical Features

```
Num_col = [col for col in df.columns if df[col].dtype != "0"]  
Num_col
```

```
['Age', 'Height', 'Weight', 'Duration', 'Heart_Rate', 'Body_Temp',  
'Calories']
```

```
df[Num_col].shape
```

```
(15000, 7)
```

```
Numerical = df[Num_col]  
Numerical.head()
```

| | Age | Height | Weight | Duration | Heart_Rate | Body_Temp | Calories |
|---|-----|--------|--------|----------|------------|-----------|----------|
| 0 | 68 | 190.0 | 94.0 | 29.0 | 105.0 | 40.8 | 231.0 |
| 1 | 20 | 166.0 | 60.0 | 14.0 | 94.0 | 40.3 | 66.0 |
| 2 | 69 | 179.0 | 79.0 | 5.0 | 88.0 | 38.7 | 26.0 |
| 3 | 34 | 179.0 | 71.0 | 13.0 | 100.0 | 40.5 | 71.0 |
| 4 | 27 | 154.0 | 58.0 | 10.0 | 81.0 | 39.8 | 35.0 |

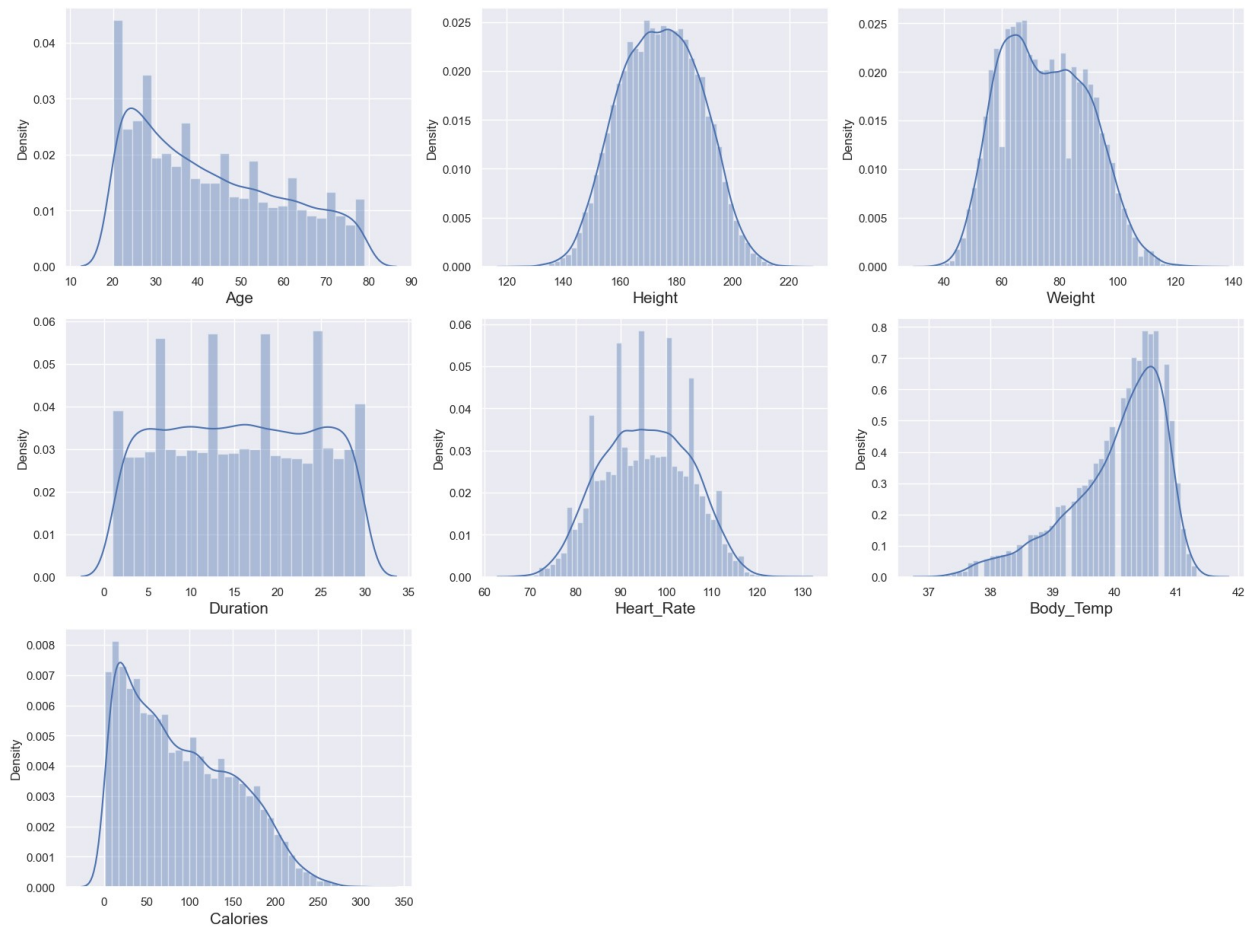
Step 4 : EDA

```
Numerical.shape
```

```
(15000, 7)
```

```
plt.figure(figsize=(20,15))  
plotnumber = 1
```

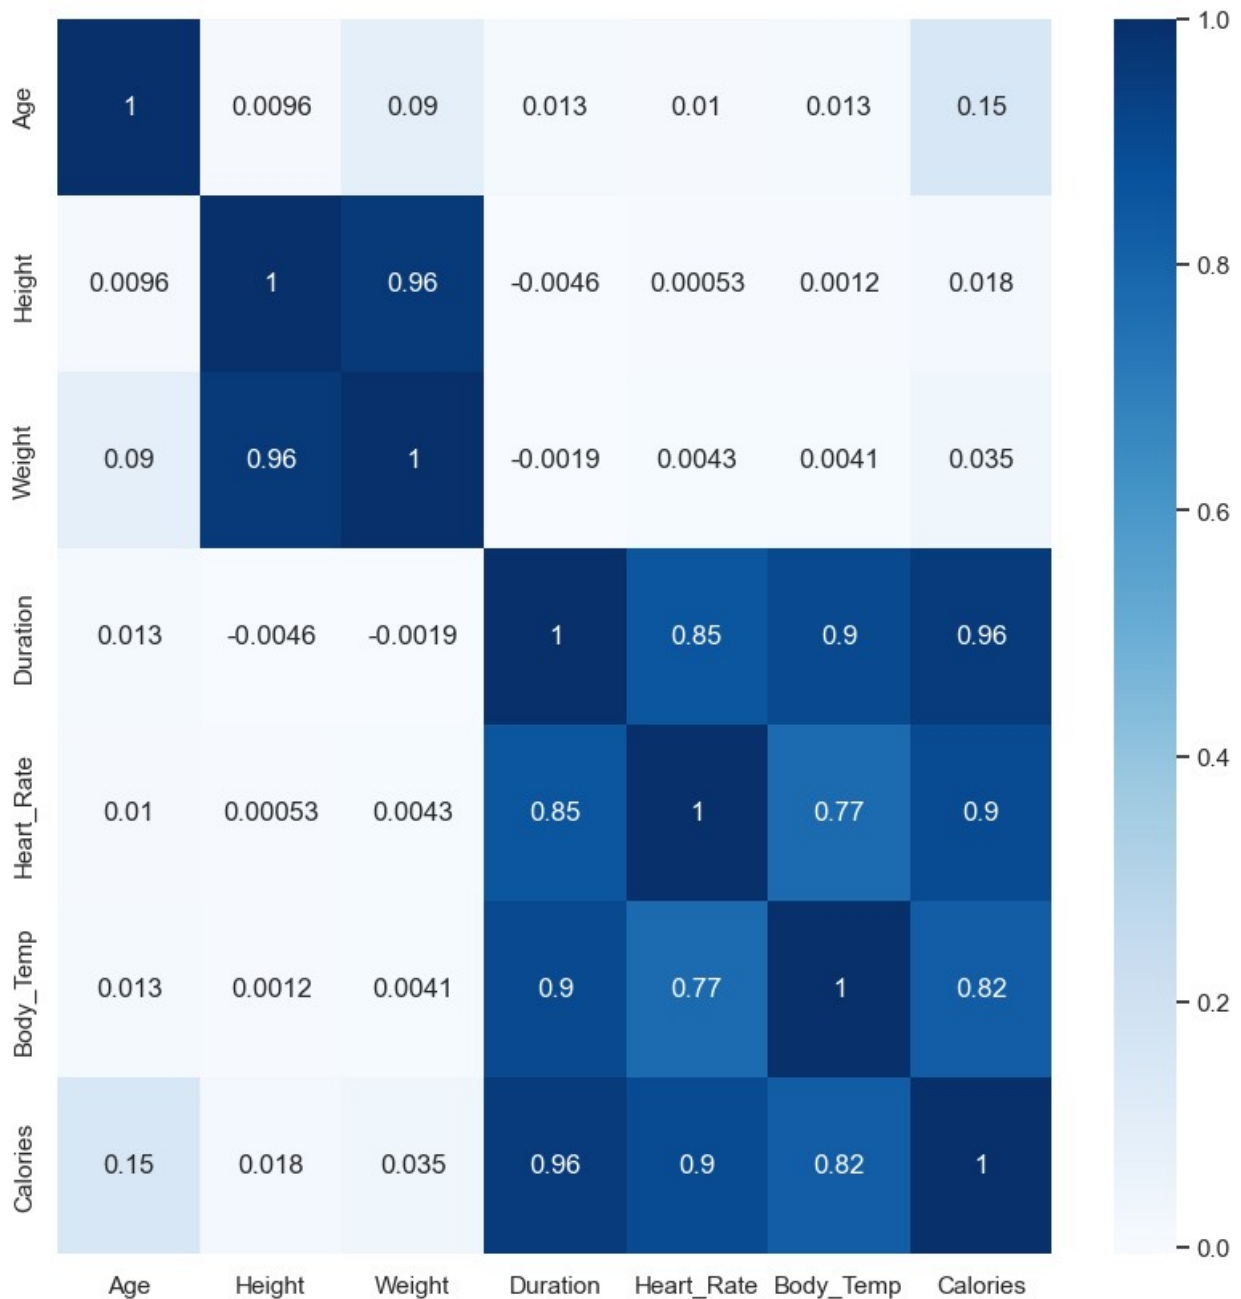
```
for column in Numerical:  
    if plotnumber <= 8:  
        ax = plt.subplot(3,3,plotnumber)  
        sns.distplot(Numerical[column])  
        plt.xlabel(column,fontsize=15)  
        plotnumber+=1  
plt.show()
```



constructing a heatmap to understand the correlation

```
plt.figure(figsize=(10,10))
sns.heatmap(Numerical.corr(), cmap='Blues',annot = True)
```

<Axes: >



Concatenate Categorical and Numerical

```
data = pd.concat([categorical, Numerical], axis=1)
```

```
data.head()
```

| | male | Age | Height | Weight | Duration | Heart_Rate | Body_Temp | Calories |
|---|-------|-----|--------|--------|----------|------------|-----------|----------|
| 0 | True | 68 | 190.0 | 94.0 | 29.0 | 105.0 | 40.8 | 231.0 |
| 1 | False | 20 | 166.0 | 60.0 | 14.0 | 94.0 | 40.3 | |

```

66.0
2  True    69   179.0   79.0     5.0     88.0     38.7
26.0
3  False   34   179.0   71.0    13.0    100.0    40.5
71.0
4  False   27   154.0   58.0    10.0     81.0    39.8
35.0


```

```

fig,ax = plt.subplots(figsize = (15,10))
sns.boxplot(data=data,width = 0.5,fliersize = 3,ax=ax)

```

<Axes: >

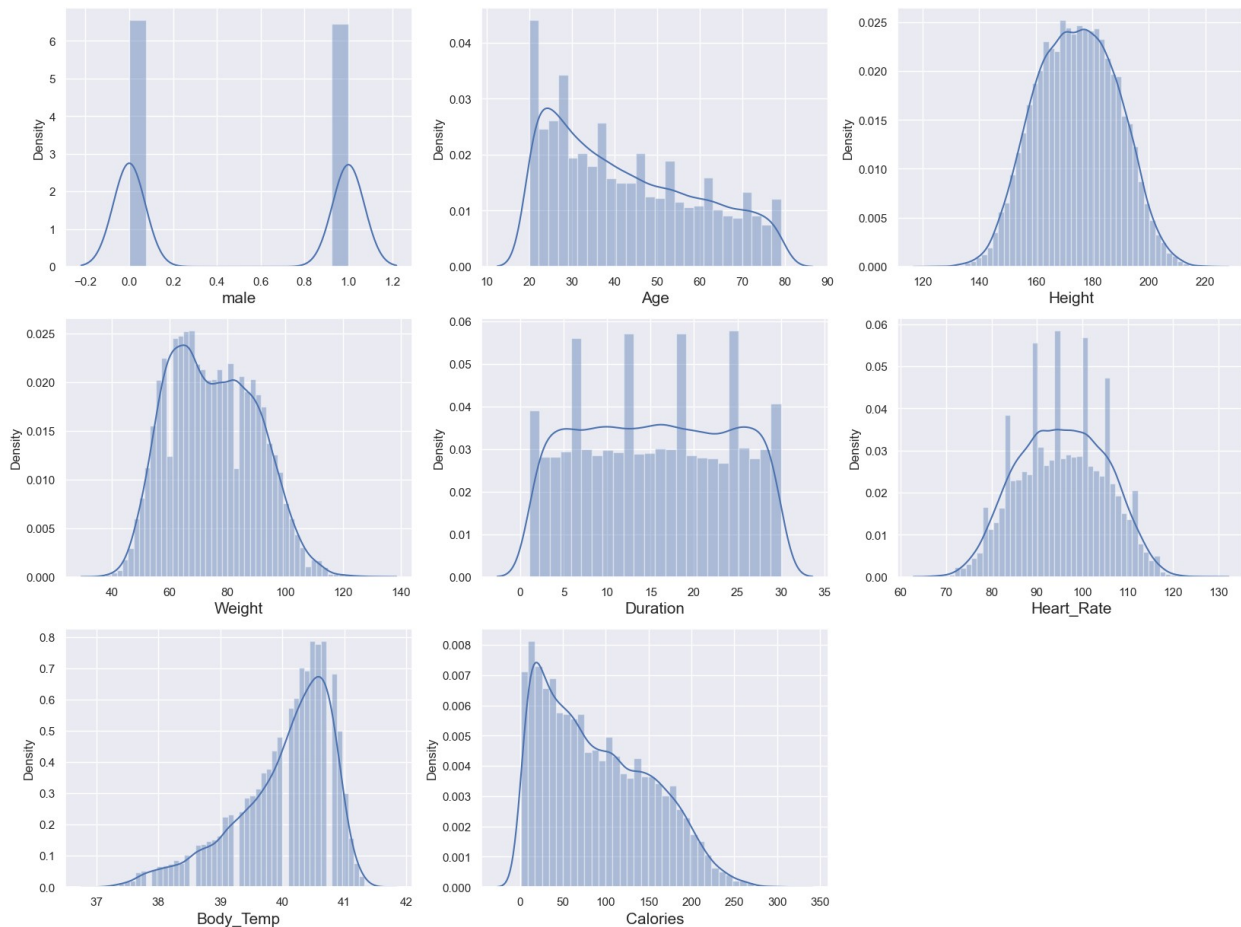


```

plt.figure(figsize=(20,15))
plotnumber = 1

for column in data:
    if plotnumber <= 8:
        ax = plt.subplot(3,3,plotnumber)
        sns.distplot(data[column])
        plt.xlabel(column,fontsize=15)
        plotnumber+=1
plt.show()

```



```
data.columns
```

```
Index(['male', 'Age', 'Height', 'Weight', 'Duration', 'Heart_Rate',  
      'Body_Temp', 'Calories'],  
      dtype='object')
```

Step 5 : Splitting Data into X and Y

```
X = data.drop(columns = ["Calories"],axis = 1)  
y = data["Calories"]
```

```
X.head()
```

| | male | Age | Height | Weight | Duration | Heart_Rate | Body_Temp |
|---|-------|-----|--------|--------|----------|------------|-----------|
| 0 | True | 68 | 190.0 | 94.0 | 29.0 | 105.0 | 40.8 |
| 1 | False | 20 | 166.0 | 60.0 | 14.0 | 94.0 | 40.3 |
| 2 | True | 69 | 179.0 | 79.0 | 5.0 | 88.0 | 38.7 |
| 3 | False | 34 | 179.0 | 71.0 | 13.0 | 100.0 | 40.5 |
| 4 | False | 27 | 154.0 | 58.0 | 10.0 | 81.0 | 39.8 |

```

y.head()
0    231.0
1     66.0
2     26.0
3     71.0
4     35.0
Name: Calories, dtype: float64

# Split the Data
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size =
0.2,random_state=1)

```

step 6 : Build a Model

```

#from sklearn import metrics
def predict(ml_model):
    model=ml_model.fit(X_train,y_train)
    print('Score : {}'.format(model.score(X_train,y_train)))
    y_prediction=model.predict(X_test)
    print('predictions are: \n {}'.format(y_prediction))
    print('\n')

    r2_score=metrics.r2_score(y_test,y_prediction)
    print('r2 score: {}'.format(r2_score))

    print('MAE:',metrics.mean_absolute_error(y_test,y_prediction))
    print('MSE:',metrics.mean_squared_error(y_test,y_prediction))

print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test,y_prediction))
)

sns.distplot(y_test-y_prediction)

```

XGB Regressor

```


regression = predict(XGBRegressor())
regression

Score : 0.9995380557081355
predictions are:
[197.06581  70.867226 196.99498  ...  29.043041 104.09284  14.61472
]

r2 score: 0.9986863132331905
MAE: 1.5521575984954834

```

MSE: 5.2744122853837005
RMSE: 2.2966088664340956



Linear Regression

```
predict(LinearRegression())
```

Score : 0.9675925554735781

predictions are:


```
[198.81182363  80.43555305 194.40940033 ... 22.14745631 118.63504926  
-11.98134672]
```

r2 score: 0.9655977245826503

MAE: 8.479071745987948

MSE: 138.12408611460904

RMSE: 11.752620393538159



DecisionTree Regression

```
predict(DecisionTreeRegressor())
```

Score : 1.0

predictions are:


```
[194.  75. 206. ...  30. 109.  13.]
```

r2 score: 0.9922832124228775

MAE: 3.544

MSE: 30.982666666666667

RMSE: 5.56207565898586



RandomForest Regression

```
predict(RandomForestRegressor())
```

Score : 0.9996811600697983

predictions are:


```
[197.68  66.23 196.22 ...  27.31 111.5   13.96]
```

r2 score: 0.9976866738104277

MAE: 1.8274800000000002

MSE: 9.287934066666667

RMSE: 3.047611206611937



```
# Logistic Regression
```

```
predict(LogisticRegression())
```

```
Score : 0.053583333333333333
```

```
predictions are:
```

```
[181.  60. 198. ...  43. 110.   3.]
```

```
r2 score: 0.836138812847468
```

```
MAE: 15.460333333333333
```

```
MSE: 657.8976666666666
```

```
RMSE: 25.6495159148602
```